

Carnegie Mellon

School of Computer Science

Deep Reinforcement Learning and Control

Multi-armed bandits

Spring 2021, CMU 10-403

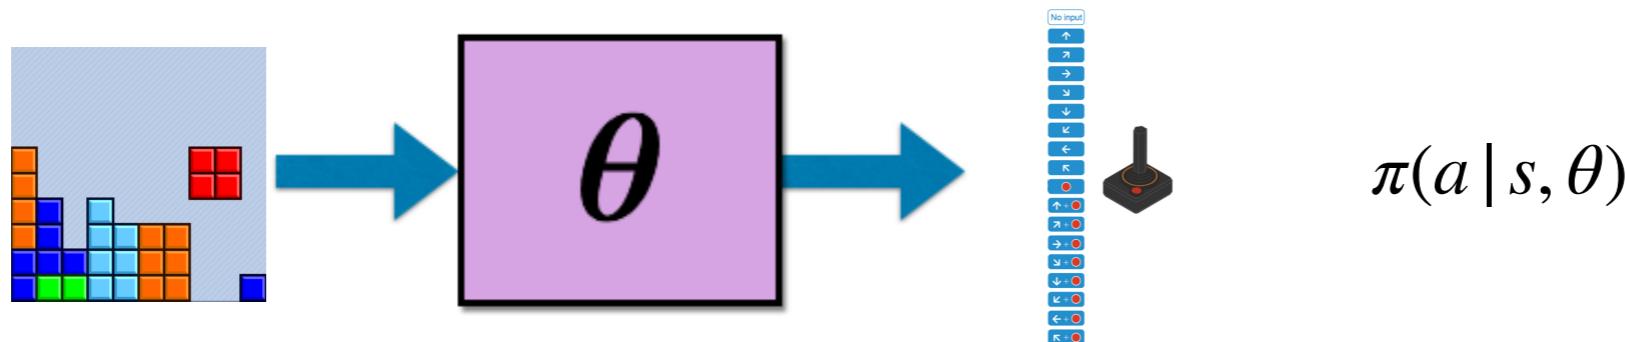
Katerina Fragkiadaki



Used Materials

- **Disclaimer:** Some material and slides for this lecture were borrowed from Rich Sutton's lecture on multi-armed bandits.

LL: Reinforcement Learning via Evolution



Given an initial state distribution $\mu_0(s_0)$, estimate parameters θ of a policy π_θ so that, the trajectories τ sampled from this policy have maximum returns, i.e., sum of rewards $R(\tau)$.

$$\max_{\theta} . \quad U(\theta) = \mathbb{E}_{\tau \sim \pi_\theta} [R(\tau) | \pi_\theta, \mu_0(s_0)]$$

τ : trajectory, a sequence of state, action, rewards, a game fragment or a full game:

$$\tau : s_0, a_0, r_0, s_1, a_1, r_1, s_2, a_2, r_2, \dots, s_T, a_T, r_T$$

$R(\tau)$: reward of a trajectory: (discounted) sum of the rewards of the individual state/actions

$$R(\tau) = \sum_{t=1}^T r_t$$

LL: Evolutionary methods for policy search

$$\max_{\theta} . \quad U(\theta) = \mathbb{E}_{\tau \sim \pi_\theta} [R(\tau) | \pi_\theta, \mu_0(s_0)]$$

General algorithm:

Initialize a population of parameter vectors (genotypes)

1. *Make random perturbations (mutations) to each parameter vector*
2. *Evaluate the perturbed parameter vector (fitness)*
3. *Keep the perturbed vector if the result improves (selection)*
4. *GOTO 1*

Simple and biologically plausible...

LL: Natural Evolutionary Strategies (NES)

- In CEM and CMA-ES, we have been selecting the best (elite) parameter offsprings.
- NES considers **every** offspring.

Algorithm 1: Evolutionary Strategies

1. **Input:** Learning rate α , noise standard deviation σ , initial policy parameters θ_0
2. **for** $t = 0, 1, 2, \dots$ **do**
3. Sample $\epsilon_1, \dots, \epsilon_n \sim \mathcal{N}(0, \mathbf{I}_d)$
4. Compute returns $F_i = F(\mu_t + \sigma\epsilon_i)$ for $i = 1, 2, \dots, n$
5. Set $\mu_{t+1} \leftarrow \mu_t + \alpha \frac{1}{n\sigma} \sum_{i=1}^n F_i \epsilon_i$
6. **end for**

This lecture - Motivation

Learning to act in a non-sequential (single action) setups:

- Each action results in an **immediate** reward.
- We want to choose actions that maximize our immediate reward in expectation.
 - Q: Why in expectation?
 - A: Because rewards are not deterministic.
- For example, displaying an advertisement can generate different click rates in different days. Actions: the advertisements to be displayed, Rewards: the user click rate. We want to pick the advertisement that maximizes the click rate on average

Multi-Armed Bandits

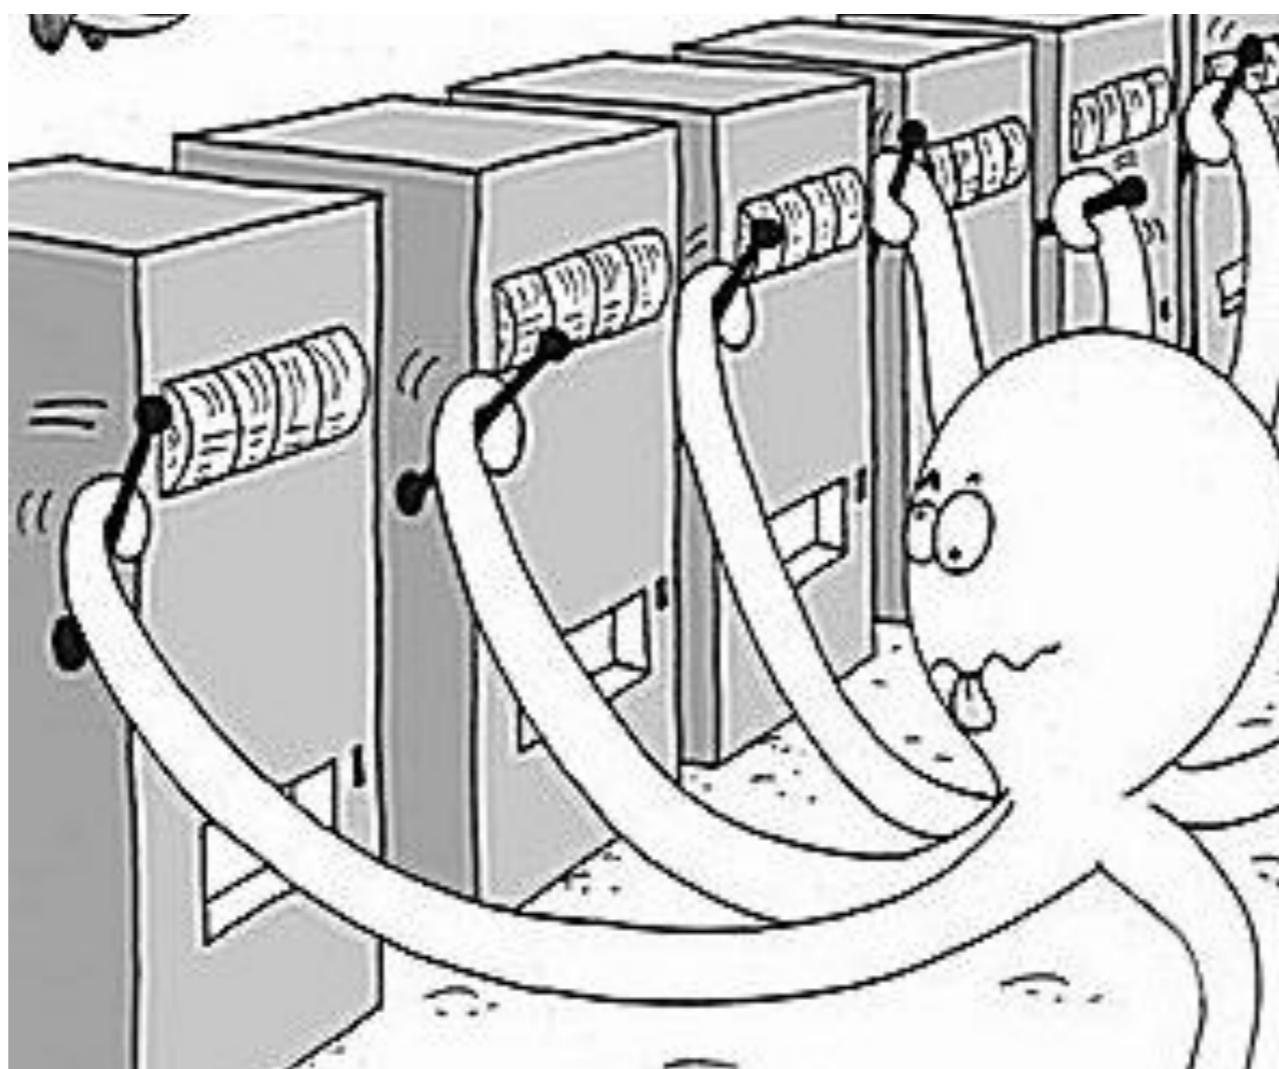
One-armed bandit= Slot machine (English slang)



source: infoslotmachine.com

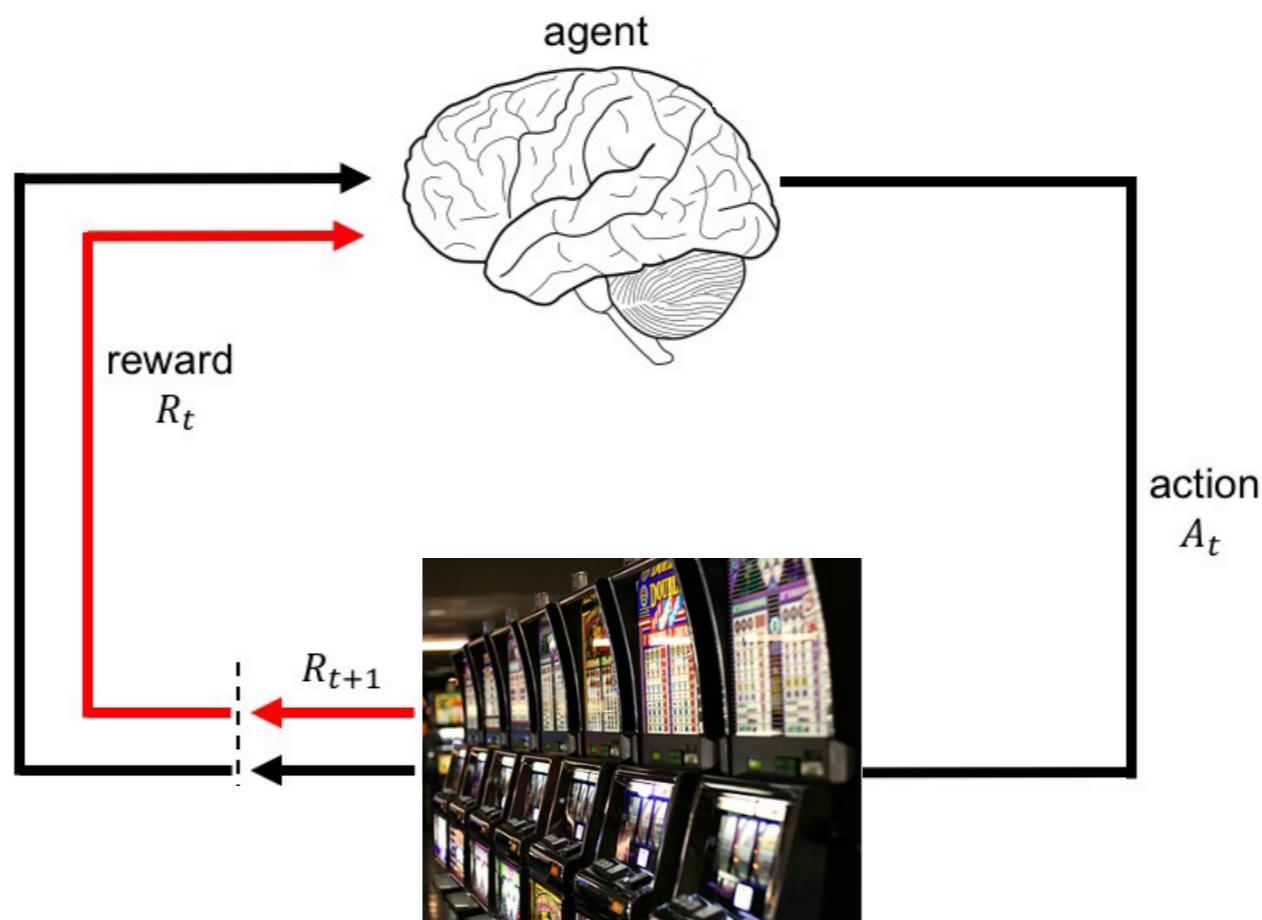
Multi-Armed Bandits

Multi-Armed bandit = Multiple Slot Machine



source: Microsoft Research

Multi-Armed Bandits



$$A_t, R_{t+1}, A_{t+1}, R_{t+2}, A_{t+2}, A_{t+3}, R_{t+3}, \dots$$

The state does not change! (a.k.a. stateless)

Multi-Armed Bandit Problem

At each timestep t the agent chooses one of the K arms and plays it.

The k th arm produces reward $r_{k,t}$ when played at timestep t .

The rewards $r_{k,t}$ are drawn from a probability distribution \mathcal{P}_k with mean μ_k .

The agent does not know neither the full arm reward distributions neither their means.



source: Pandey et al.'s slide

Agent's Objective: Maximize cumulative rewards (over an infinite horizon).

This is equivalent to **finding the arm with the highest mean reward. Why?**

Multi-Armed Bandit Problem

At each timestep t the agent chooses one of the K arms and plays it.

The k th arm produces reward $r_{k,t}$ when played at timestep t .

The rewards $r_{k,t}$ are drawn from a probability distribution \mathcal{P}_k with mean μ_k .

The agent does not know neither the full arm reward distributions neither their means.



source: Pandey et al.'s slide

Definition: The **action-value** for action α (here arm k) is **its mean reward**:

$$q_*(a) \doteq \mathbb{E}[R_t | A_t = a]$$

The Exploration/Exploitation Dilemma

- Suppose you form estimates

$$Q_t(a) \approx q_*(a), \quad \forall a \quad \textit{action-value estimates}$$

The Exploration/Exploitation Dilemma

- Suppose you form estimates

$$Q_t(a) \approx q_*(a), \quad \forall a \quad \text{action-value estimates}$$

- Define the *greedy action* at time t as

$$A_t^* \doteq \arg \max_a Q_t(a)$$

The Exploration/Exploitation Dilemma

- Suppose you form estimates

$$Q_t(a) \approx q_*(a), \quad \forall a \quad \text{action-value estimates}$$

- Define the *greedy action* at time t as

$$A_t^* \doteq \arg \max_a Q_t(a)$$

- If $A_t = A_t^*$ then you are *exploiting*

If $A_t \neq A_t^*$ then you are *exploring*

The Exploration/Exploitation Dilemma

- Suppose you form estimates

$$Q_t(a) \approx q_*(a), \quad \forall a \quad \text{action-value estimates}$$

- Define the *greedy action* at time t as

$$A_t^* \doteq \arg \max_a Q_t(a)$$

- If $A_t = A_t^*$ then you are *exploiting*
If $A_t \neq A_t^*$ then you are *exploring*
- You can't do both, but you need to do both

The Exploration/Exploitation Dilemma

- Suppose you form estimates

$$Q_t(a) \approx q_*(a), \quad \forall a \quad \text{action-value estimates}$$

- Define the *greedy action* at time t as

$$A_t^* \doteq \arg \max_a Q_t(a)$$

- If $A_t = A_t^*$ then you are *exploiting*
If $A_t \neq A_t^*$ then you are *exploring*
- You can't do both, but you need to do both
- You can never stop exploring, but maybe you should explore less with time.

Exploration vs Exploitation Dilemma

- Online decision-making involves a fundamental choice:
 - **Exploitation:** Make the best decision given current information
 - **Exploration:** Gather more information
- The best long-term strategy may involve **short-term sacrifices**
- Gather enough information to make the best overall decisions

Exploration vs Exploitation Dilemma

- Online decision-making involves a fundamental choice:
 - **Exploitation:** Make the best decision given current information
 - **Exploration:** Gather more information
- The best long-term strategy may involve **short-term sacrifices**
- Gather enough information to make the best overall decisions
- The exploration/exploitation dilemma is not a problem encountered in computational RL or deep RL: It is a fundamental problem in decision making of any intelligent agent.

Exploration vs. Exploitation Dilemma

- Restaurant Selection
 - **Exploitation:** Go to your favorite restaurant
 - **Exploration:** Try a new restaurant
- Oil Drilling
 - **Exploitation:** Drill at the best known location
 - **Exploration:** Drill at a new location
- Game Playing
 - **Exploitation:** Play the move you believe is best
 - **Exploration:** Play an experimental move

Example: Bernoulli Bandits

Recall: The **Bernoulli distribution** is the discrete probability distribution of a random variable which takes the value 1 with probability p and the value 0 with probability $q=1-p$, that is, the probability distribution of any single experiment that asks a yes–no question.

- Each action (arm when played) results in success or failure, **rewards are binary**.
- Mean reward for each arm represents the probability of success.
- Action (arm) $k \in \{1\dots K\}$ produces a success with probability $\theta_k \in [0,1]$.



$$\theta_1$$

win 0.6
of time



$$\theta_2$$

win 0.4
of time

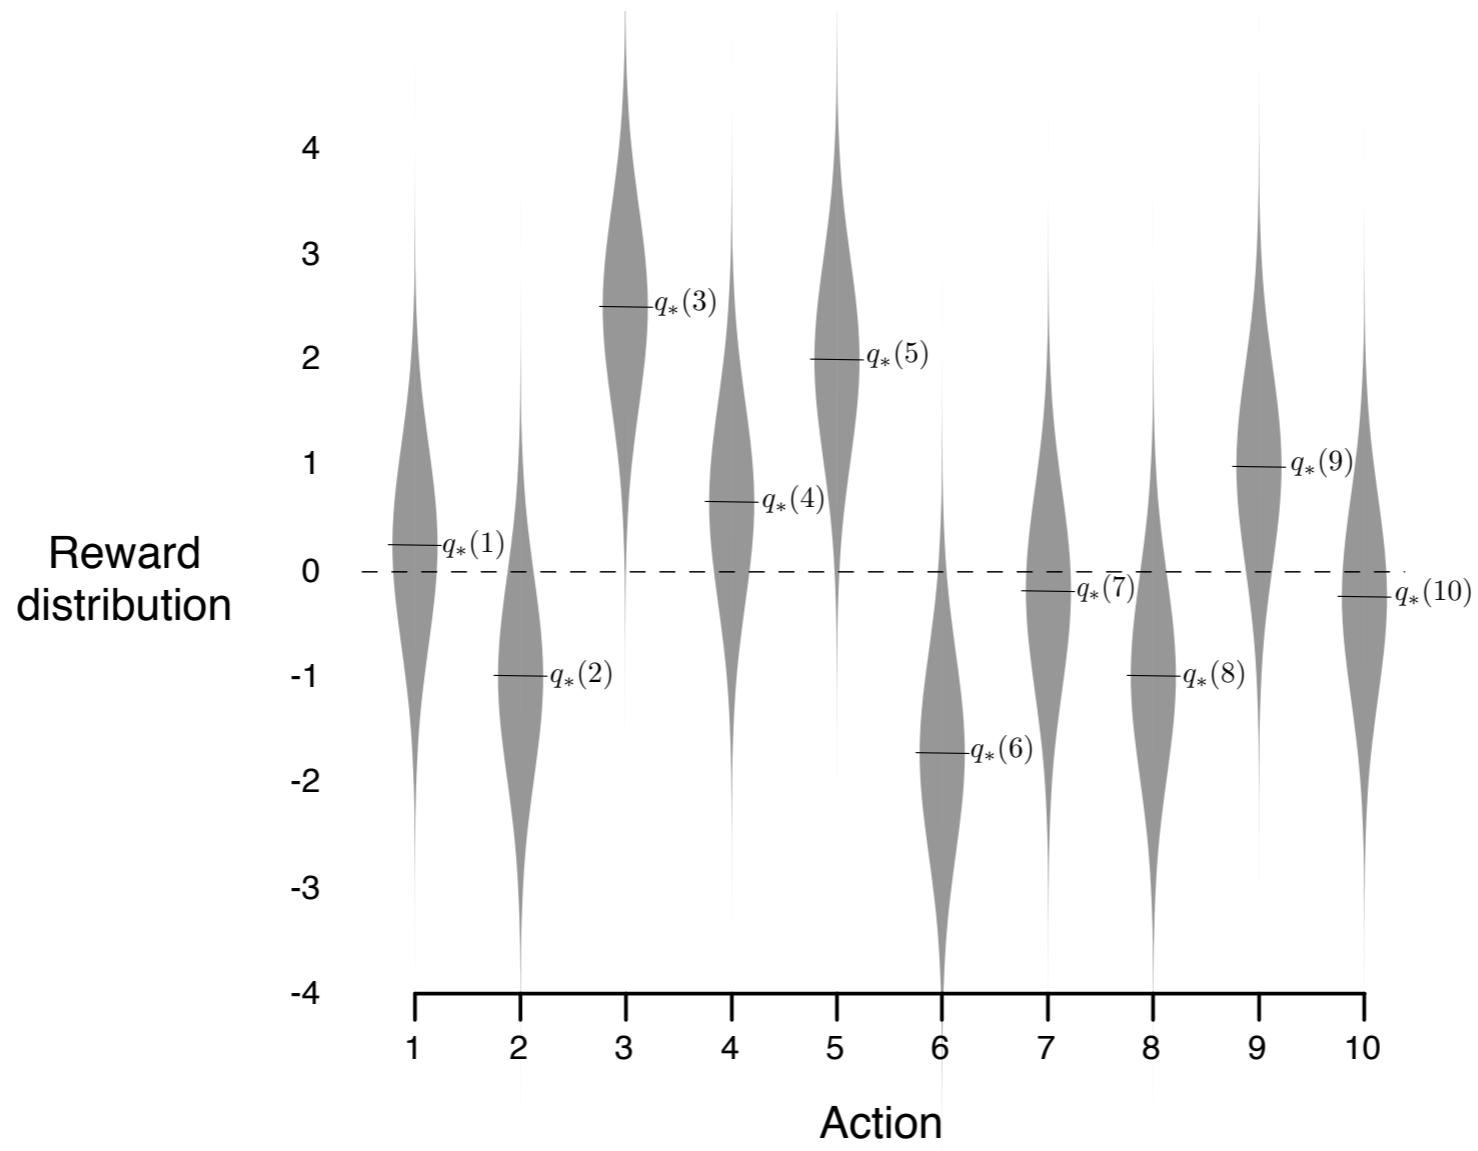


$$\theta_3$$

win 0.45
of time

Example: Gaussian Bandits

- Each action (arm when played) results in a **real number**.
- Action (arm) $k \in \{1 \dots K\}$ produces on average reward equal to the mean of its Gaussian distribution.



Real world motivation: content presentation

We have two variations of content of a webpage, A and B, and we want to decide which one to display to engage more users.

- Two arm bandits: each arm corresponds to a content variation shown to users (not necessarily the same user).
- Mean rewards: the total percentage of users that would click on each invitation

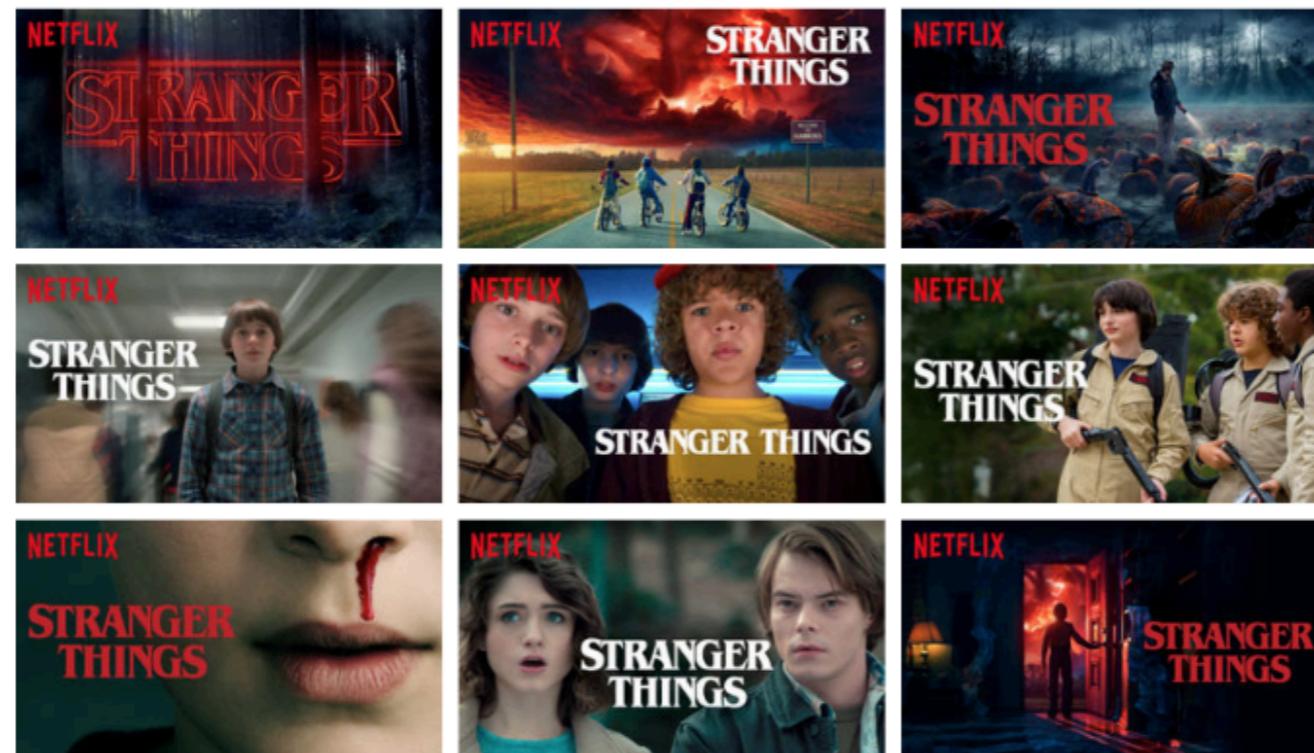
The screenshot shows a promotional offer for a dinner with President Barack Obama. The main headline is "DINNER WITH BARACK" with the subtext "Your chance to meet the President". Below the headline is a large image of President Obama smiling. Text overlay on the image reads "DINNER WITH BARACK" and "YOU'RE INVITED. WE'LL COVER YOUR AIRFARE." At the bottom of the page, there is a detailed legal disclaimer about the promotion's rules and restrictions, followed by standard footer links for "OBAMA BIDEN", "Privacy Policy", and "Terms of Service".

The screenshot shows the same promotional offer for a dinner with President Barack Obama. The main headline is "DINNER WITH BARACK" with the subtext "Your chance to meet the President". Below the headline is a large image of the First Family (President Obama, First Lady Michelle Obama, and their daughters) sitting around a table. Text overlay on the image reads "DINNER WITH BARACK" and "You're invited. We'll cover your airfare." At the bottom of the page, there is a detailed legal disclaimer about the promotion's rules and restrictions, followed by standard footer links for "OBAMA BIDEN", "Privacy Policy", and "Terms of Service".

Real world motivation: NETFLIX artwork

For a particular movie, we want to decide what image to show (to all the NETFLIX users)

- Actions: uploading one of the K images to a user's home screen
- Mean rewards: the percentage of users that clicked and watched (quality engagement, not clickbait)



Netflix Artwork

Regret

- The **action-value** is the mean reward for action a ,

$$q_*(a) \doteq \mathbb{E}[R_t | A_t = a], \quad \forall a \in \{1, \dots, k\}$$

- The **optimal value** is

$$\nu_* = q(a^*) = \max_{a \in \mathcal{A}} q_*(a)$$

- The **regret** is the opportunity loss for one step

$$I_t = \mathbb{E}[\nu_* - q_*(a_t)] \quad \text{reward} = -\text{regret}$$

- The **total regret** is the total opportunity loss

$$L_T = \mathbb{E} \left[\sum_{t=1}^T \nu_* - q_*(a_t) \right]$$

- Maximize cumulative reward = minimize total regret

Regret

- The **count** $N_t(a)$: the number of times that action a has been selected prior to time t
- The **gap** Δ_a is the difference in value between action a and optimal action a_* : $\Delta_a = v_* - q_*(a)$
- Regret is a function of gaps and the counts

$$\begin{aligned} L_T &= \mathbb{E} \left[\sum_{t=1}^T v_* - q_*(a_t) \right] \\ &= \sum_{a \in \mathcal{A}} \mathbb{E}[N_t(a)](v_* - q_*(a)) \\ &= \sum_{a \in \mathcal{A}} \mathbb{E}[N_t(a)]\Delta_a \end{aligned}$$

Forming Action-Value Estimates

- To simplify notation, let us focus on one action
 - We consider only its rewards, and its estimate after $n+1$ rewards:

$$Q_n \doteq \frac{R_1 + R_2 + \cdots + R_{n-1}}{n - 1}$$

- How can we do this incrementally (without storing all the rewards)?
- Could store a running sum and count (and divide), or equivalently:

$$Q_{n+1} = Q_n + \frac{1}{n} [R_n - Q_n]$$

- This is a standard form for learning/update rules:

$$\text{NewEstimate} \leftarrow \text{OldEstimate} + \text{StepSize} [\text{Target} - \text{OldEstimate}]$$

Forming Action-Value Estimates

- To simplify notation, let us focus on one action
 - We consider only its rewards, and its estimate after $n+1$ rewards:

$$Q_n \doteq \frac{R_1 + R_2 + \cdots + R_{n-1}}{n - 1}$$

- How can we do this incrementally (without storing all the rewards)?
- Could store a running sum and count (and divide), or equivalently:

$$Q_{n+1} = Q_n + \frac{1}{n} [R_n - Q_n]$$

- This is a standard form for learning/update rules: error

$$\text{NewEstimate} \leftarrow \text{OldEstimate} + \text{StepSize} \begin{bmatrix} \text{Target} - \text{OldEstimate} \end{bmatrix}$$

Derivation of incremental update

$$Q_n \doteq \frac{R_1 + R_2 + \cdots + R_{n-1}}{n - 1}$$

$$\begin{aligned} Q_{n+1} &= \frac{1}{n} \sum_{i=1}^n R_i \\ &= \frac{1}{n} \left(R_n + \sum_{i=1}^{n-1} R_i \right) \\ &= \frac{1}{n} \left(R_n + (n-1) \frac{1}{n-1} \sum_{i=1}^{n-1} R_i \right) \\ &= \frac{1}{n} \left(R_n + (n-1)Q_n \right) \\ &= \frac{1}{n} \left(R_n + nQ_n - Q_n \right) \\ &= Q_n + \frac{1}{n} [R_n - Q_n], \end{aligned}$$

Non-stationary bandits

- Suppose the true action values change slowly over time
 - then we say that the problem is *nonstationary*
- In this case, sample averages are not a good idea
 - Why?

Non-stationary bandits

- Suppose the true action values change slowly over time
 - then we say that the problem is *nonstationary*
- In this case, sample averages are not a good idea
- Better is an “exponential, recency-weighted average”:

$$\begin{aligned} Q_{n+1} &\doteq Q_n + \alpha \left[R_n - Q_n \right] \\ &= (1 - \alpha)^n Q_1 + \sum_{i=1}^n \alpha (1 - \alpha)^{n-i} R_i, \end{aligned}$$

where $\alpha \in (0,1]$ and constant

The smaller the i , the smaller the multiplier-> forgetting earlier rewards

Action selection in multi-armed bandits

Fixed exploration period + Greedy

1. Allocate a fixed time period to exploration when you try bandits **uniformly at random**
2. Estimate mean rewards for all actions:
$$Q_t(a) = \frac{1}{N_t(a)} \sum_{i=1}^{t-1} r_i \mathbf{1}(A_i = a)$$
3. Select the action that is optimal for the estimated mean rewards, breaking ties at random:
$$a_t = \operatorname{argmax}_{a \in \mathcal{A}} Q_t(a)$$
4. GOTO 2

Fixed exploration period + Greedy

- After the fixed exploration period we have formed the following reward estimates



$$Q_t(a_1) = 0.3$$



$$Q_t(a_2) = 0.5$$



$$Q_t(a_3) = 0.1$$

Q1: Will the greedy method always pick the second action?

Q2: Can greedy lock onto a suboptimal action forever?

⇒ Greedy has linear total regret

ϵ -Greedy Action Selection

- In greedy action selection, you always exploit
- In ϵ -greedy, you are usually greedy, but with probability ϵ you instead pick an action at random (possibly the greedy action again)
- This is perhaps the simplest way to balance exploration and exploitation

ϵ -Greedy Action Selection

A simple bandit algorithm

Initialize, for $a = 1$ to k :

$$\begin{aligned} Q(a) &\leftarrow 0 \\ N(a) &\leftarrow 0 \end{aligned}$$

Repeat forever:

$$\begin{aligned} A &\leftarrow \begin{cases} \arg \max_a Q(a) & \text{with probability } 1 - \varepsilon \\ \text{a random action} & \text{with probability } \varepsilon \end{cases} \quad (\text{breaking ties randomly}) \\ R &\leftarrow \text{bandit}(A) \\ N(A) &\leftarrow N(A) + 1 \\ Q(A) &\leftarrow Q(A) + \frac{1}{N(A)} [R - Q(A)] \end{aligned}$$

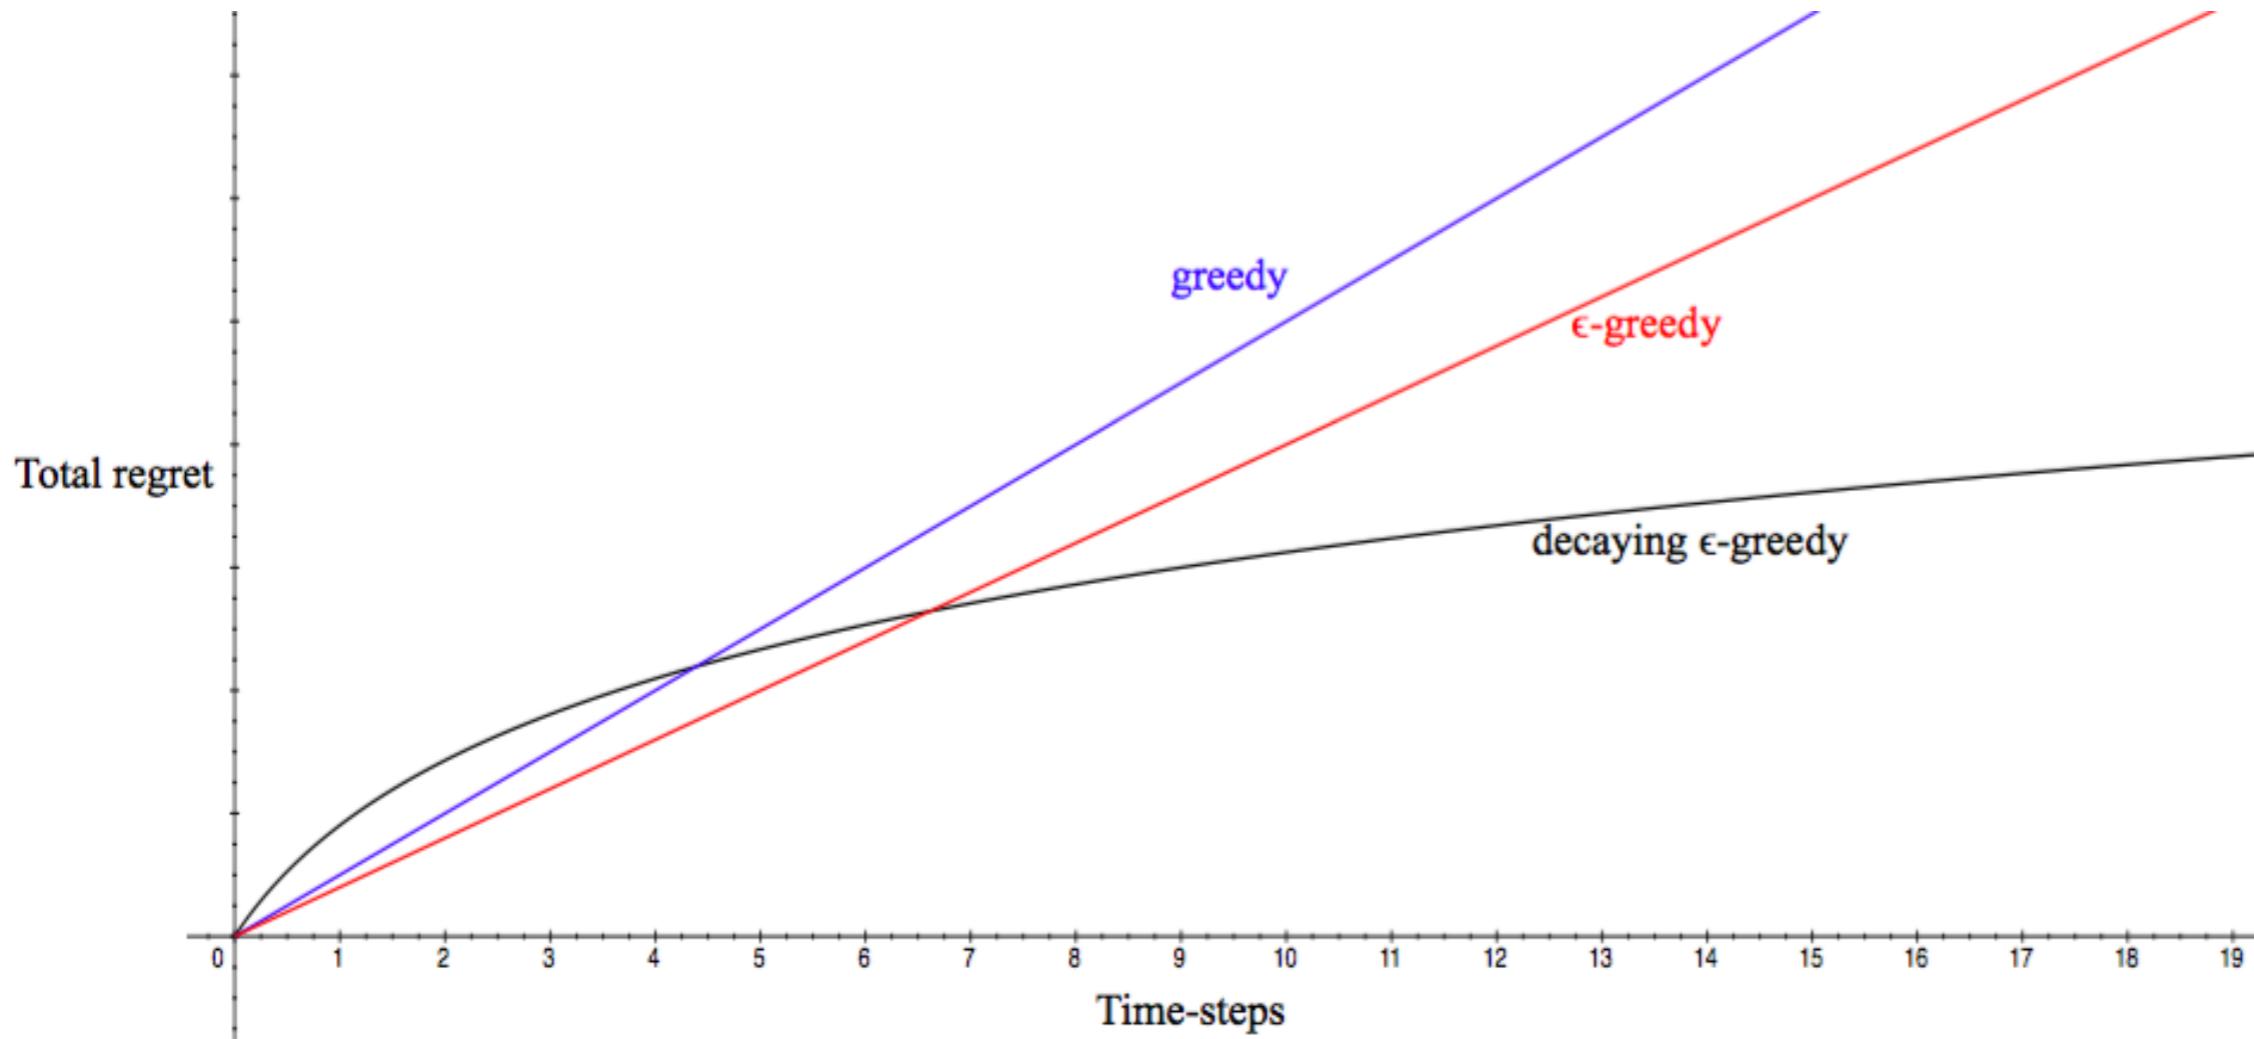
ϵ -Greedy Algorithm

- The ϵ -greedy algorithm continues to explore forever
 - With probability $1 - \epsilon$ select $a_t = \operatorname{argmax}_{a \in \mathcal{A}} Q_t(a)$
 - With probability ϵ select a random action
- Constant ϵ ensures minimum regret

$$I_t \geq \frac{\epsilon}{|\mathcal{A}|} \sum_{a \in \mathcal{A}} \Delta_a$$

- $\Rightarrow \epsilon$ -greedy has linear total regret

Counting Regret

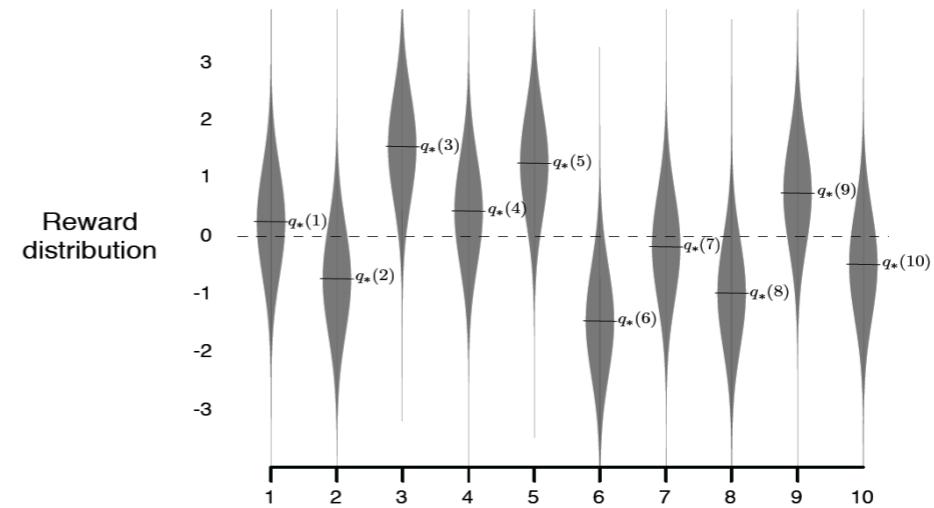


- If an algorithm forever explores it will have linear total regret
- If an algorithm never explores it will have linear total regret

Average reward for three algorithms

We sample 10 arm bandits instantiations:

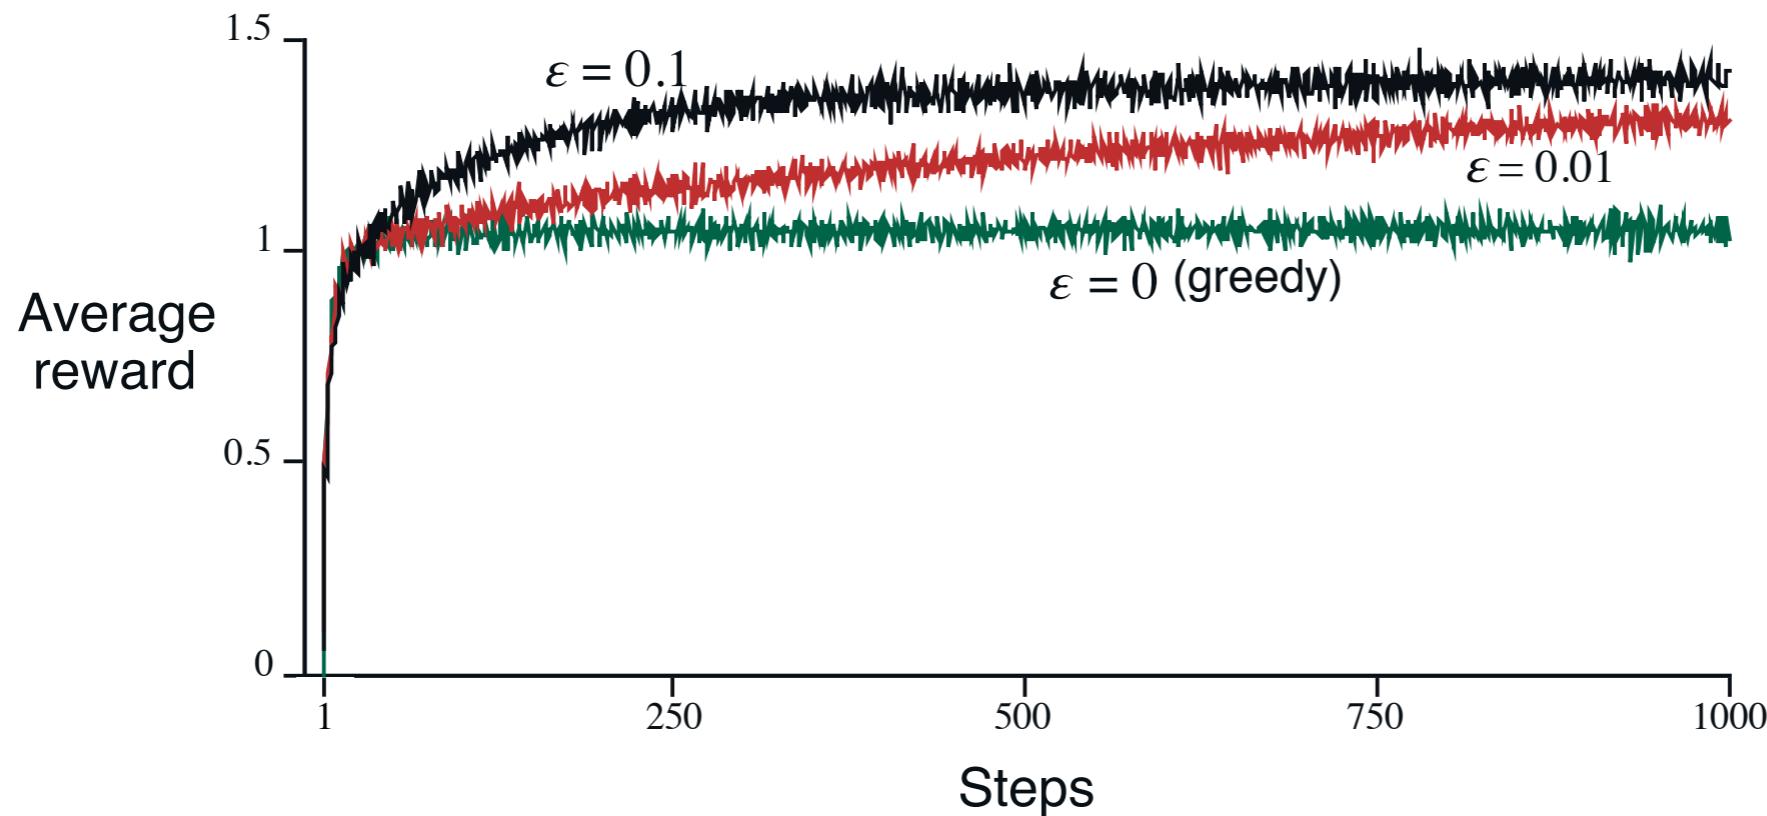
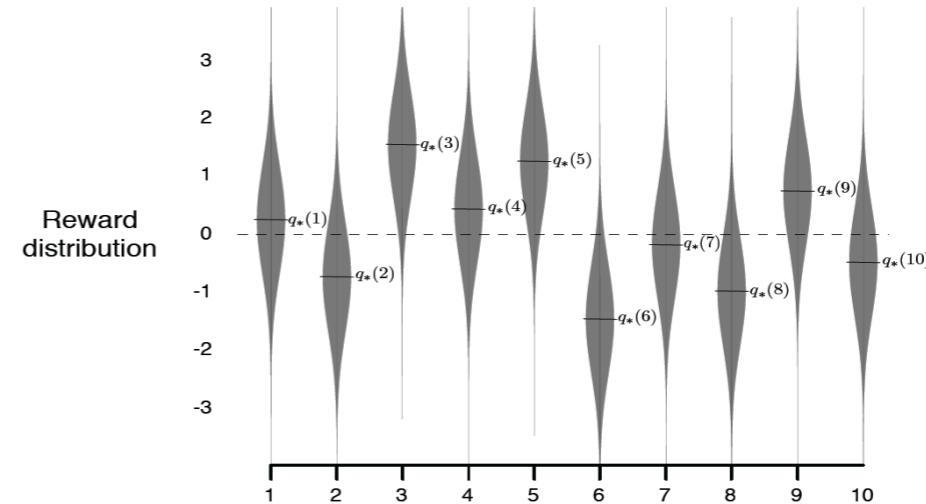
$$q_*(a) \sim \mathcal{N}(0, 1)$$
$$R_t \sim \mathcal{N}(q_*(a), 1)$$



Average reward for three algorithms

We sample 10 arm bandits instantiations:

$$q_*(a) \sim \mathcal{N}(0, 1)$$
$$R_t \sim \mathcal{N}(q_*(a), 1)$$

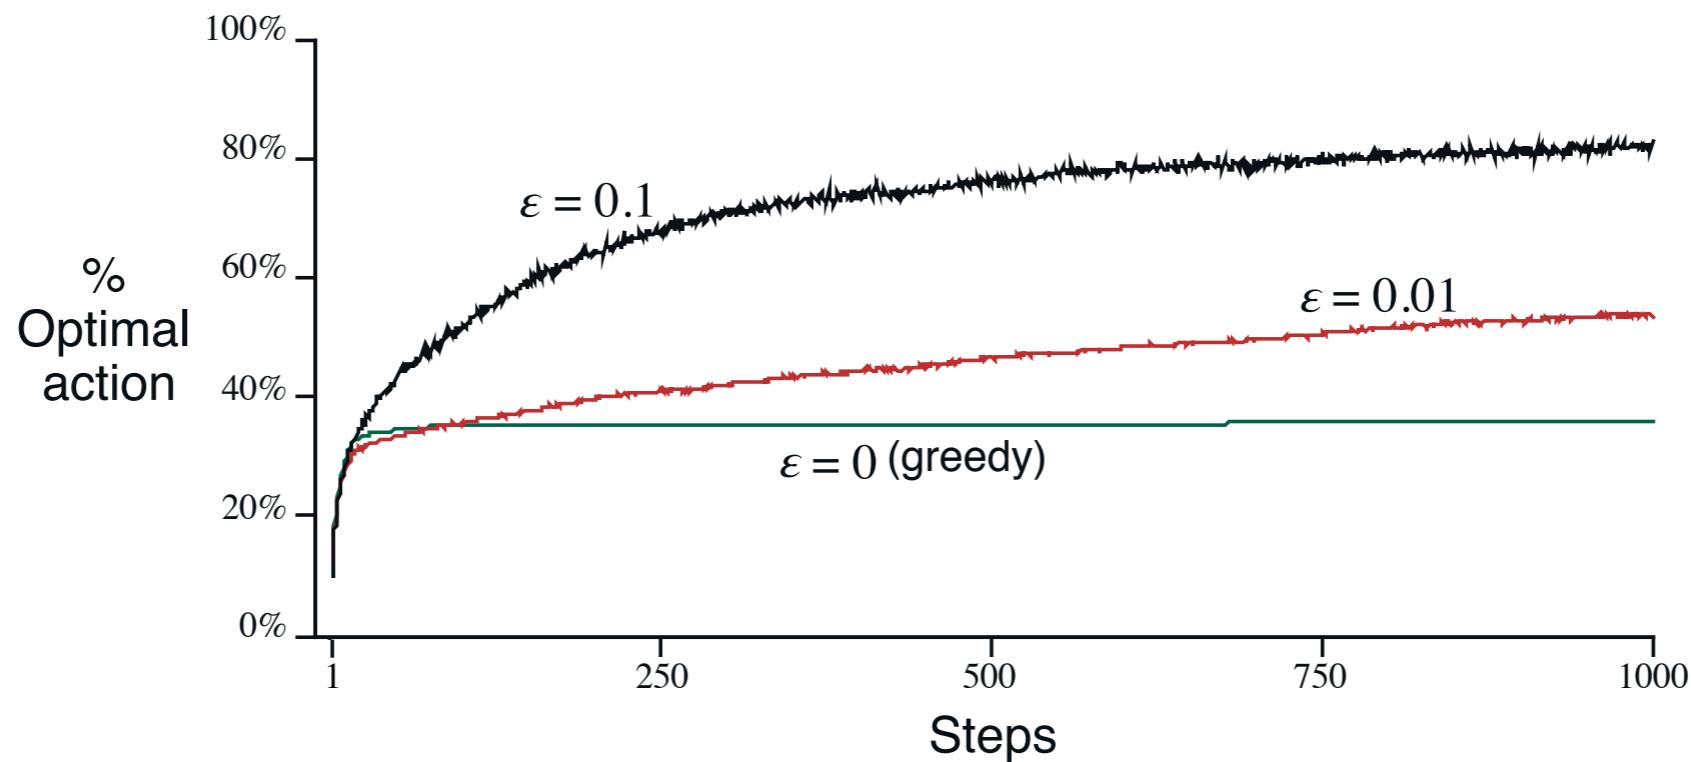
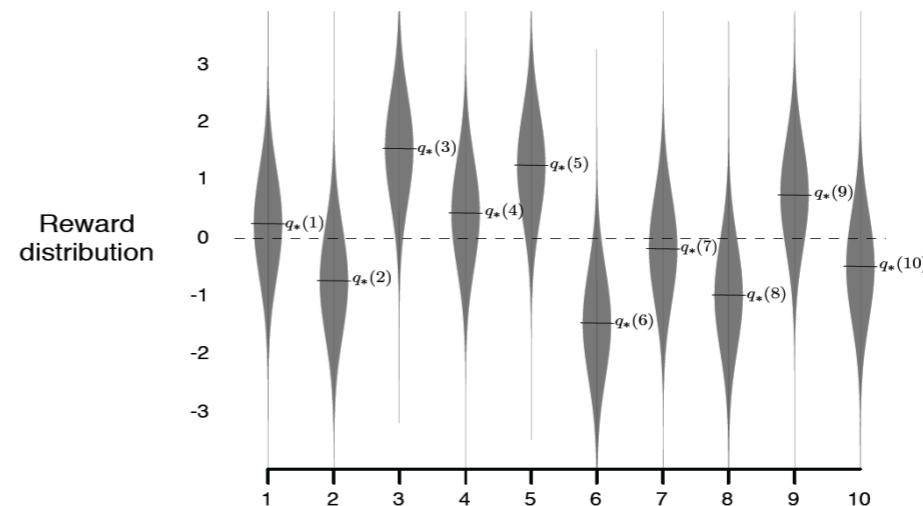


Q: In the limit (after infinite number of steps), which method will result in the largest average reward?

Optimal action for three algorithms

We sample 10 arm bandits instantiations:

$$q_*(a) \sim \mathcal{N}(0, 1)$$
$$R_t \sim \mathcal{N}(q_*(a), 1)$$

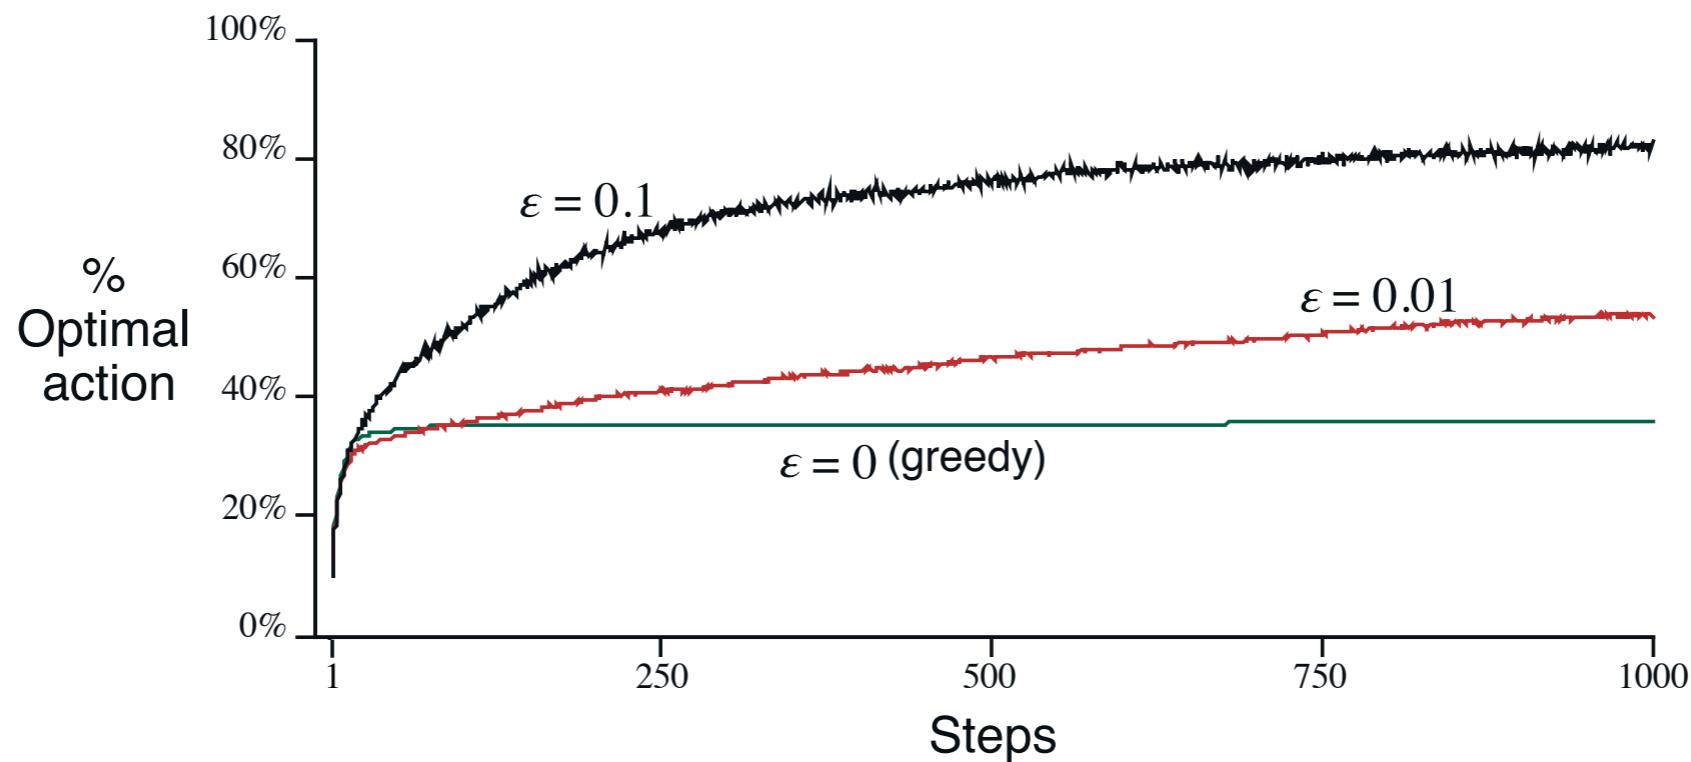
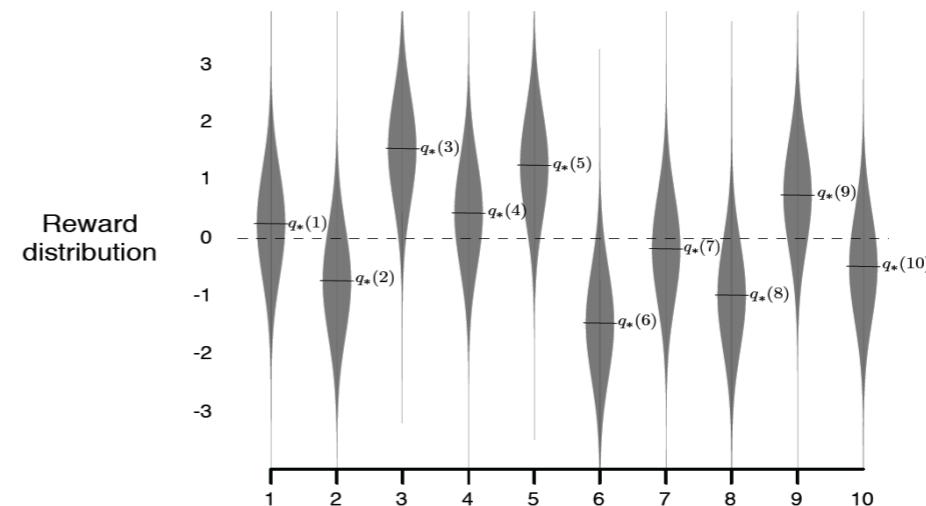


Q: Which method will find the optimal action in the limit?

Optimal action for three algorithms

We sample 10 arm bandits instantiations:

$$q_*(a) \sim \mathcal{N}(0, 1)$$
$$R_t \sim \mathcal{N}(q_*(a), 1)$$



Q: Does the performance of those methods depend on the initialization of the action value estimates?

Optimistic Initialization

- Simple and practical ideas: initialise $Q(a)$ to a high value
- Update action value by incremental Monte-Carlo evaluation
- Starting with $N(a) > 0$,

$$Q_t(a_t) = Q_{t-1}(a_t) + \frac{1}{N_t(a_t)} (r_t - Q_{t-1}(a_t))$$

just an incremental estimate of sample mean, including one 'hallucinated' initial optimistic value

- Encourages systematic exploration early on
- But optimistic greedy can still lock onto a suboptimal action if rewards are stochastic.

Optimistic Initial Values

We initialize with the following reward estimates for Bernoulli bandits



$$Q_t(a_1) = 1$$



$$Q_t(a_2) = 1$$



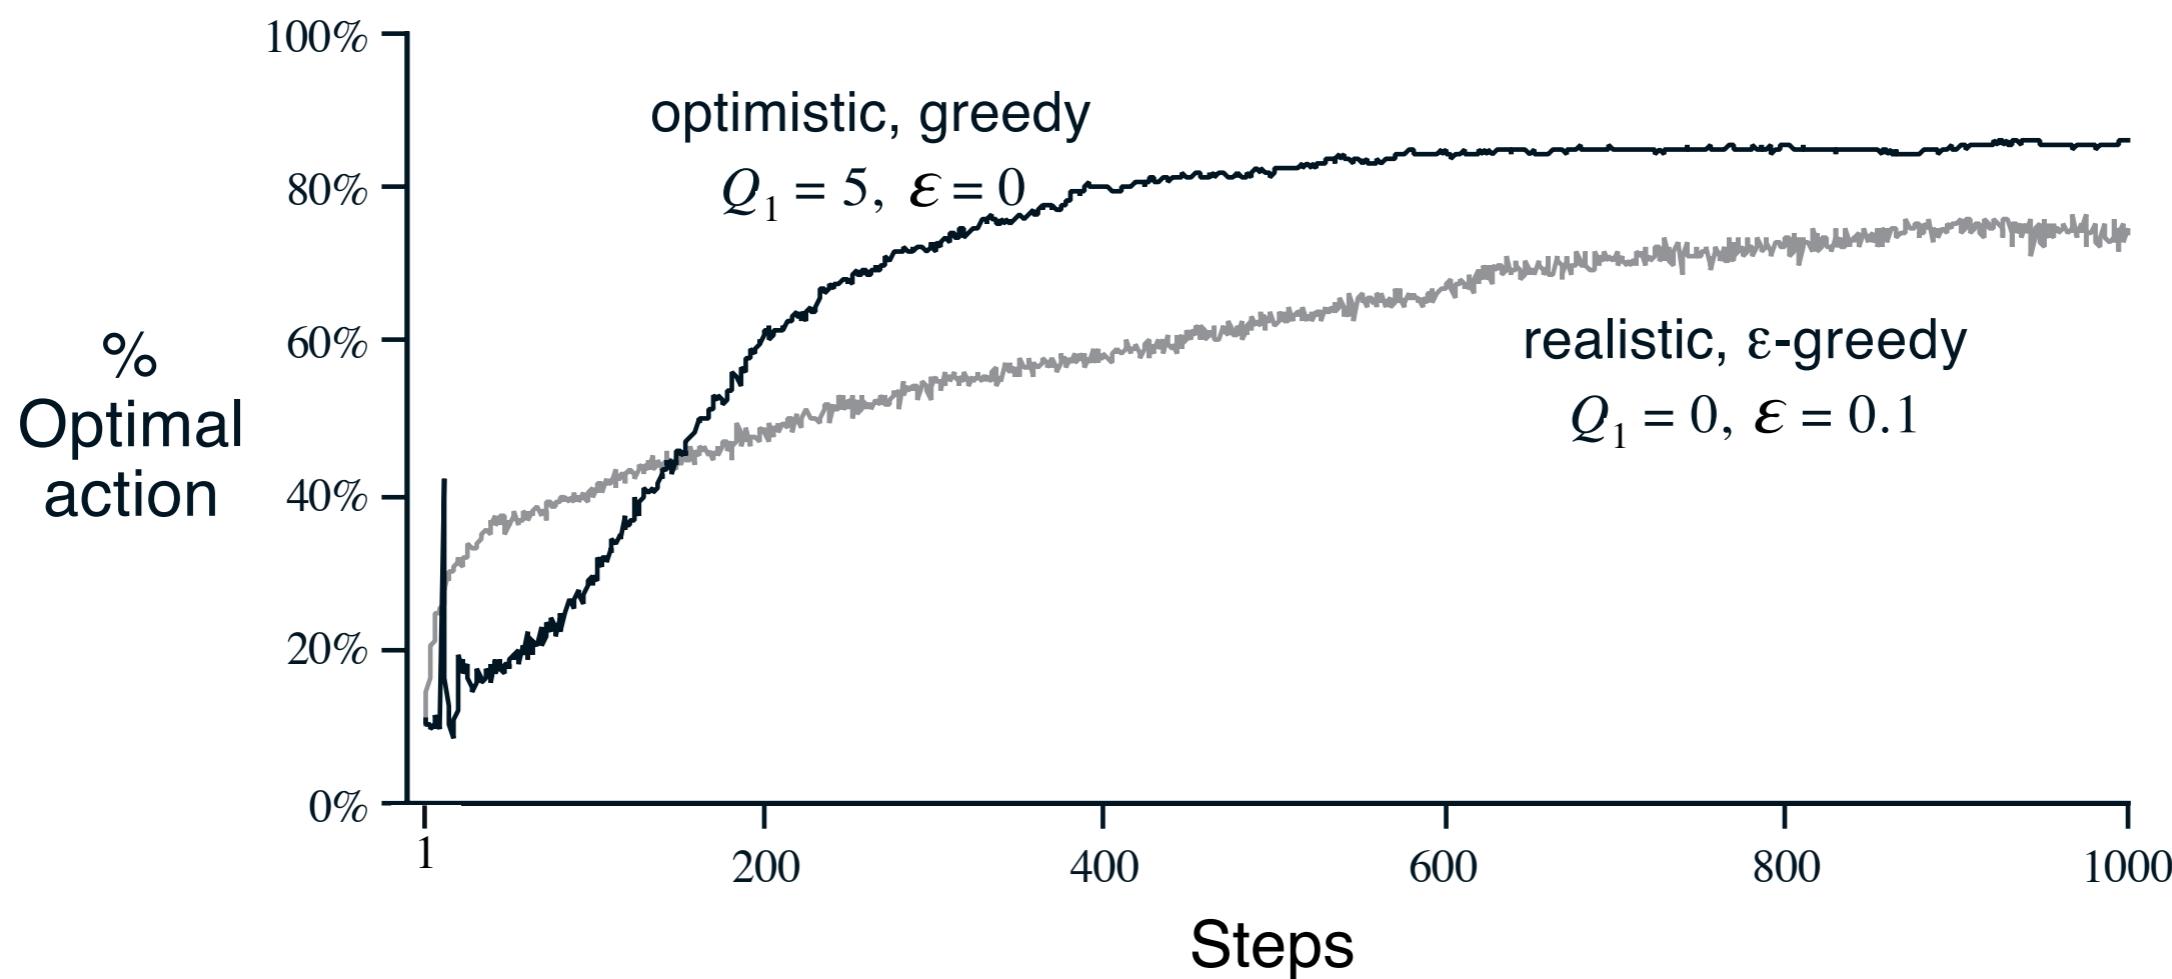
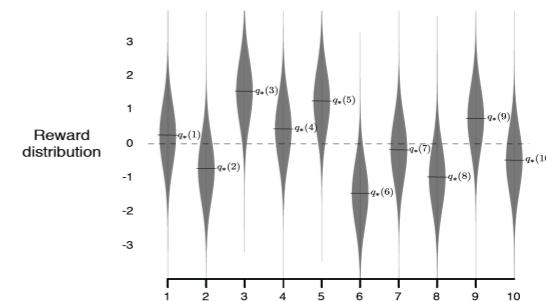
$$Q_t(a_3) = 1$$

Q: When it is possible that greedy action selection will not try out all the actions?

Optimistic Initial Values

- Suppose we initialize the action values optimistically ($Q_1(a) = 5$), e.g., on the 10-armed testbed

$$q_*(a) \sim \mathcal{N}(0, 1)$$
$$R_t \sim \mathcal{N}(q_*(a), 1)$$



Achieving sub-linear total regret

We need to reason about **uncertainty** of our action value estimates



1000 pulls,
600 wins
 $Q_t(a_1)=0.6$



1000 pulls,
400 wins
 $Q_t(a_2)=0.4$



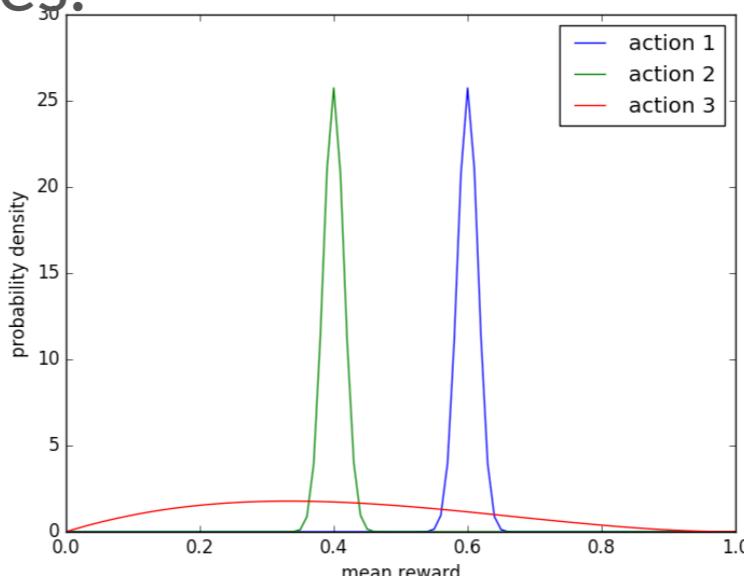
10 pulls,
4 wins
 $Q_t(a_1)=0.4$

Epsilon-greedy

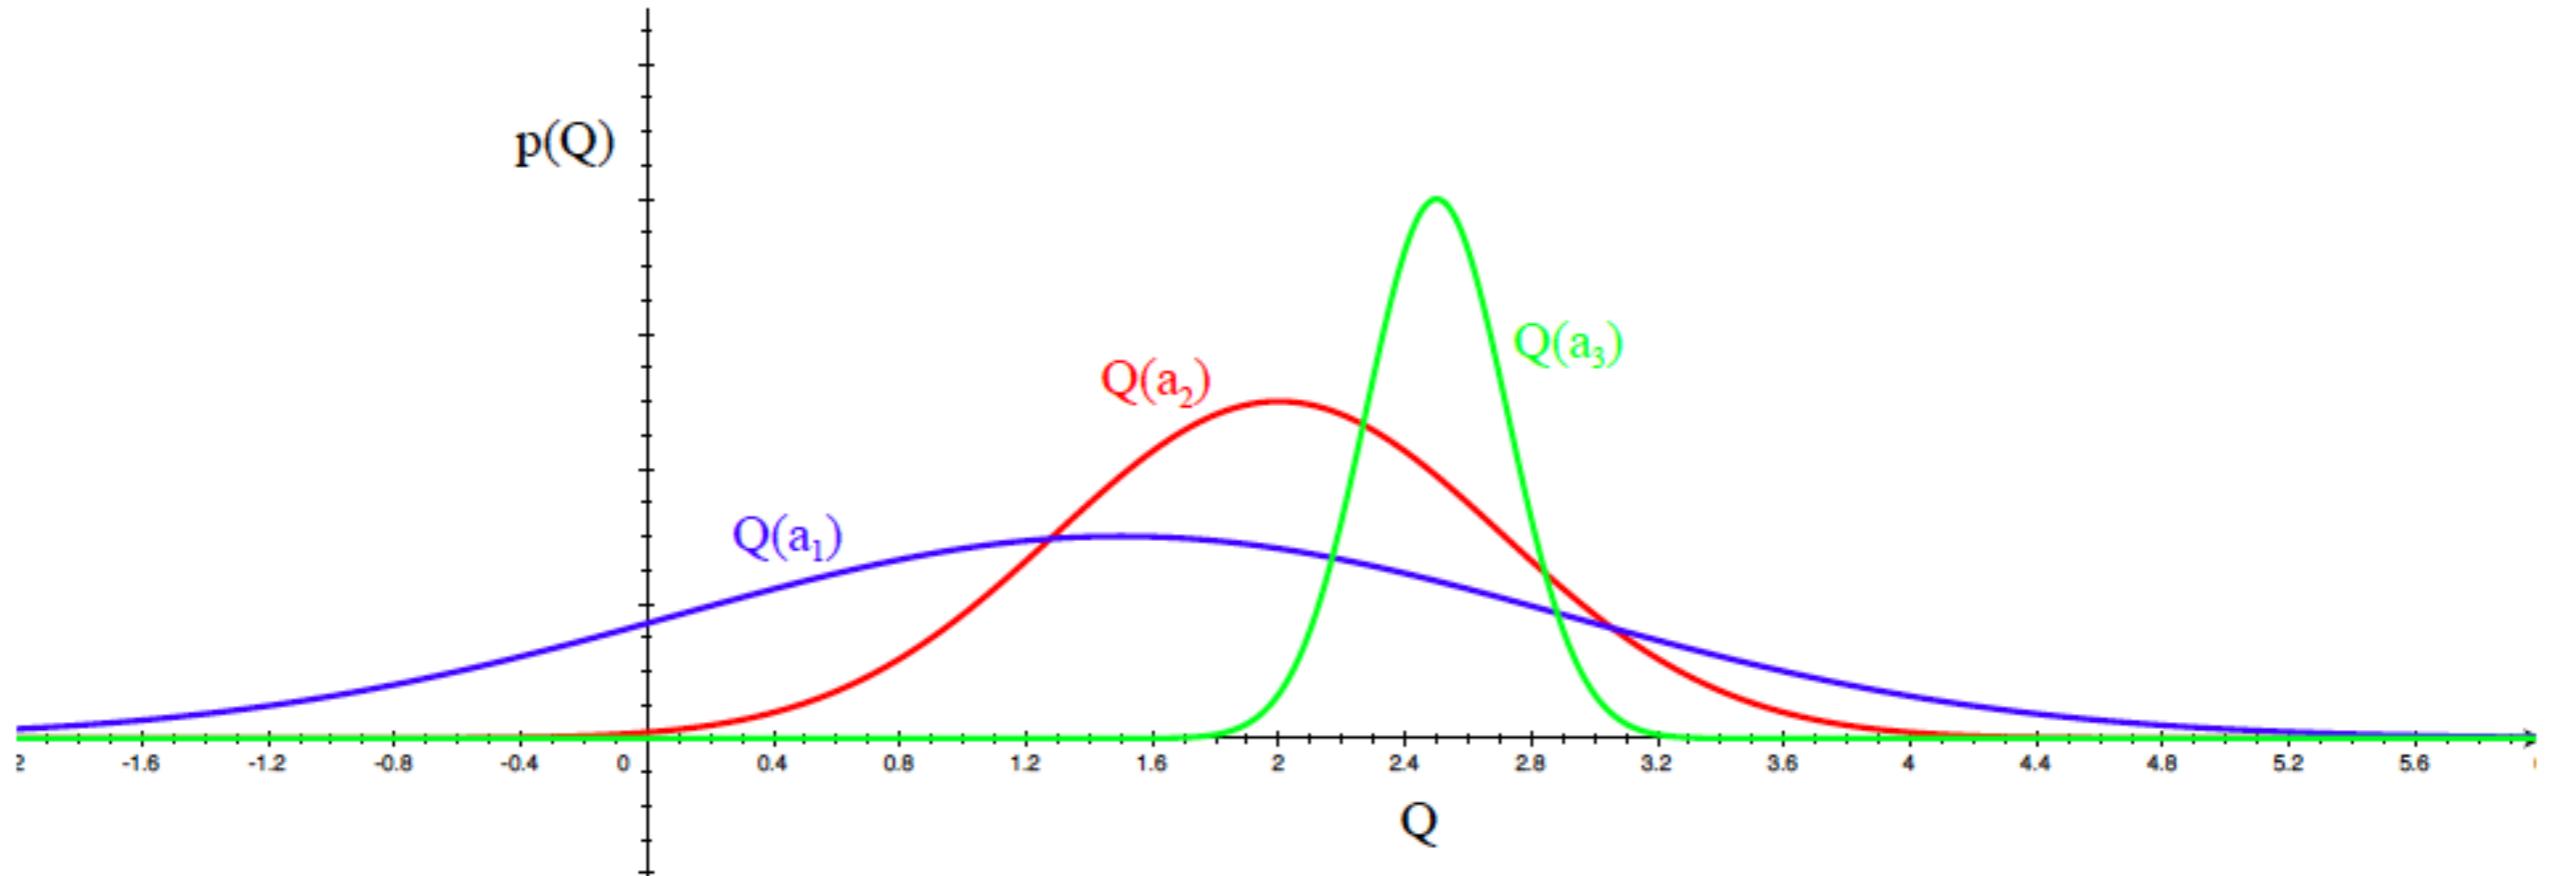
Repeat forever:
$$A \leftarrow \begin{cases} \arg \max_a Q(a) & \text{with probability } 1 - \varepsilon \\ \text{a random action} & \text{with probability } \varepsilon \end{cases}$$

 $R \leftarrow \text{bandit}(A)$
 $N(A) \leftarrow N(A) + 1$
$$Q(A) \leftarrow Q(A) + \frac{1}{N(A)} [R - Q(A)]$$

The problem with using mean estimates is that we do not reason about **uncertainty** of those estimates.

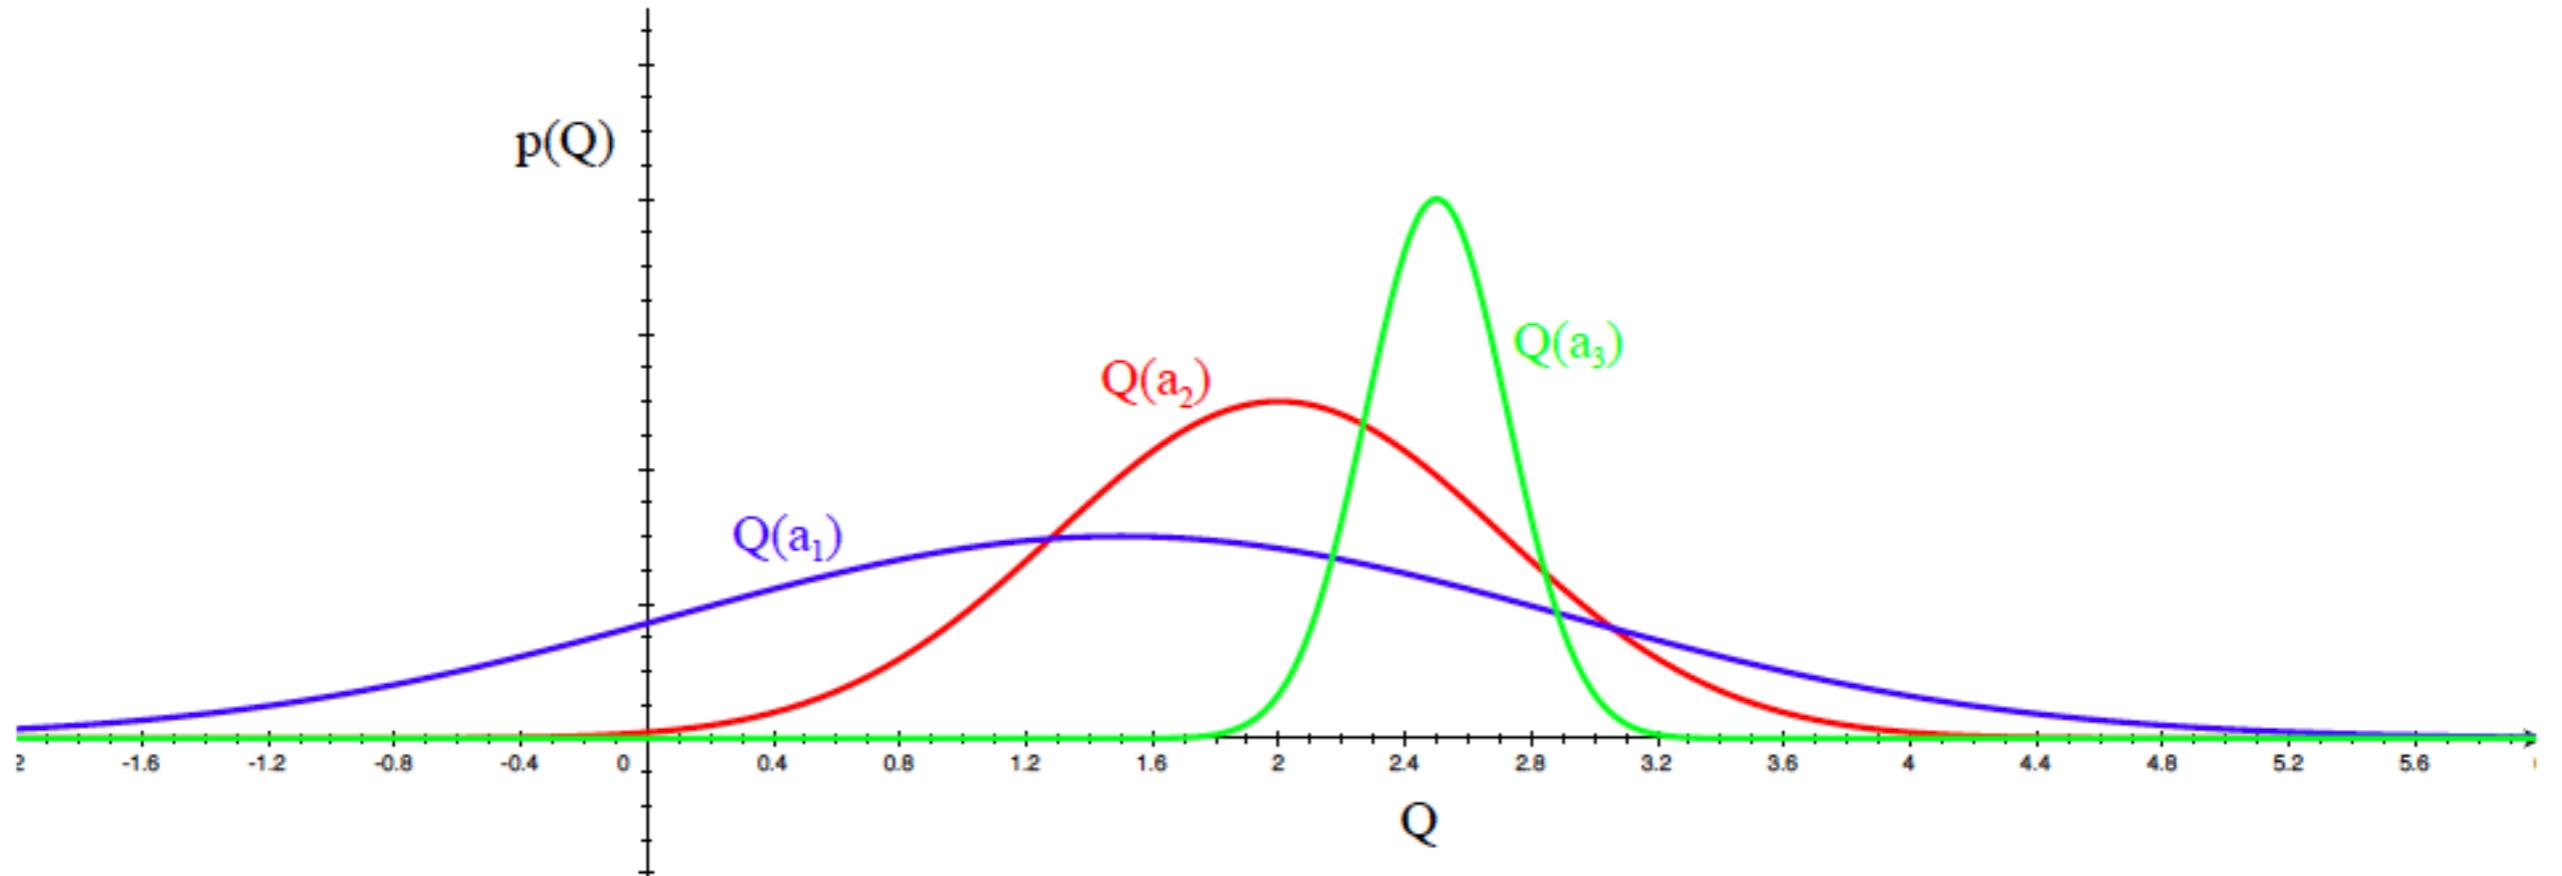


Uncertainty guides Exploration



- The more **uncertain** we are about an action-value
- The more important it is to explore that action
- It could turn out to be the best action

Uncertainty guides Exploration



- We are then less uncertain about the value
- And more likely to pick another action
- Until we converge to the best action

Upper Confidence Bounds

- Estimate an **upper confidence** $U_t(a)$ for each action value
- Such that with high probability

$$q_*(a) \leq Q_t(a) + U_t(a)$$



- This upper confidence depends on the number of times action a has been selected
 - Small $N_t(a)$ \Rightarrow large $U_t(a)$ (estimated value is uncertain)
 - Large $N_t(a)$ \Rightarrow small $U_t(a)$ (estimated value is accurate)
- Select action maximizing **Upper Confidence Bound** (UCB)

$$a_t = \operatorname{argmax}_{a \in \mathcal{A}} Q_t(a) + U_t(a)$$

Upper Confidence Bound (UCB)

- A clever way of reducing exploration over time
- Estimate an upper bound on the true action values
- Select the action with the largest (estimated) upper bound

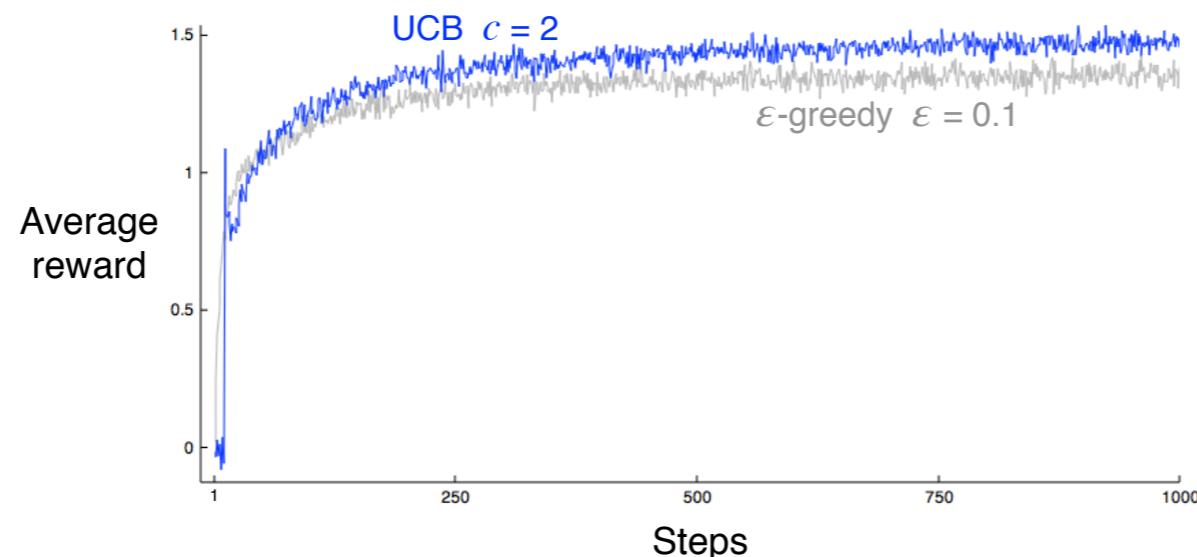
$$A_t \doteq \arg \max_a \left[Q_t(a) + c \sqrt{\frac{\log t}{N_t(a)}} \right]$$

- c is a hyper-parameter that trades-off explore/exploit
- the confidence bound grows with the total number of actions we have taken t but shrinks with the number of times we have tried this particular action $N_t(a)$. This ensures each action is tried infinitely often but still balances exploration and exploitation.

Upper Confidence Bound (UCB)

- A clever way of reducing exploration over time
- Estimate an upper bound on the true action values
- Select the action with the largest (estimated) upper bound

$$A_t \doteq \arg \max_a \left[Q_t(a) + c \sqrt{\frac{\log t}{N_t(a)}} \right]$$



UCB1 regret bound

- The regret of UCB1 at a rate of $\ln t$. After t actions it is at most:

$$\sum_{i=1}^K \frac{4 \ln t}{\Delta_i} + \left(1 + \frac{\pi^2}{3}\right) \Delta_i$$

where $\Delta_i = \mu^* - \mu_i$

Bayesian Bandits

- So far we have made no assumptions about the reward distributions.
 - In UCB we just considered some bounds on rewards
- Bayesian bandits exploit **prior knowledge of rewards**, $p[\mathcal{R}]$
- They compute **posterior** distribution of rewards $p[\mathcal{R} \mid h_t]$
$$h_t = a_1, r_1, \dots, a_{t-1}, r_{t-1}$$
- Use posterior to guide exploration: we simply sample from the posterior!

Bayesian learning for model parameters

Step 1: Given n data, $D = x_{1\dots n} = \{x_1, x_2, \dots, x_n\}$ write down the expression for likelihood:

$$p(\textcolor{red}{D} | \theta)$$

Step 2: Specify a prior: $p(\theta)$

Step 3: Compute the posterior:

$$p(\theta | D) = \frac{p(D | \theta)p(\theta)}{p(D)}$$

Thompson Sampling

Represent a distribution for the mean reward of each bandit as opposed to its mean estimate alone. At each timestep:

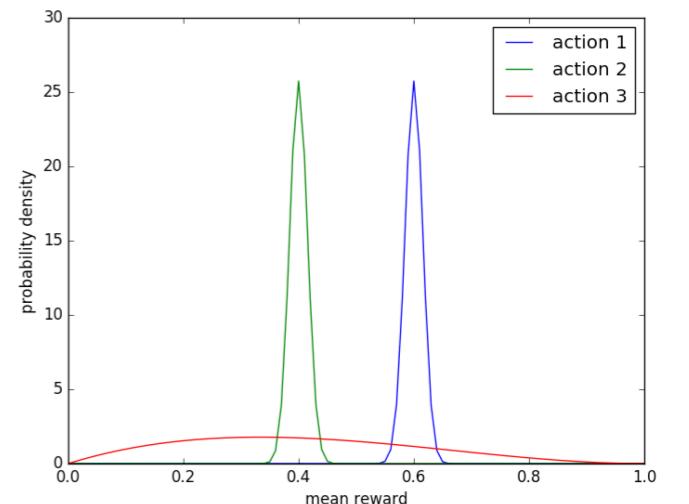
1. Sample from the reward distribution:

$$\bar{\theta}_1 \sim \hat{p}(\theta_1), \bar{\theta}_2 \sim \hat{p}(\theta_2), \dots, \bar{\theta}_k \sim \hat{p}(\theta_k)$$

2. Choose action $a = \arg \max_a \mathbb{E}_{\bar{\theta}}[r(a)]$

3. Observe the reward

4. Update the mean reward posterior distributions: $\hat{p}(\theta_1), \hat{p}(\theta_2) \dots \hat{p}(\theta_k)$



Q: why we use argmax in step 2 and we do not add any noise?

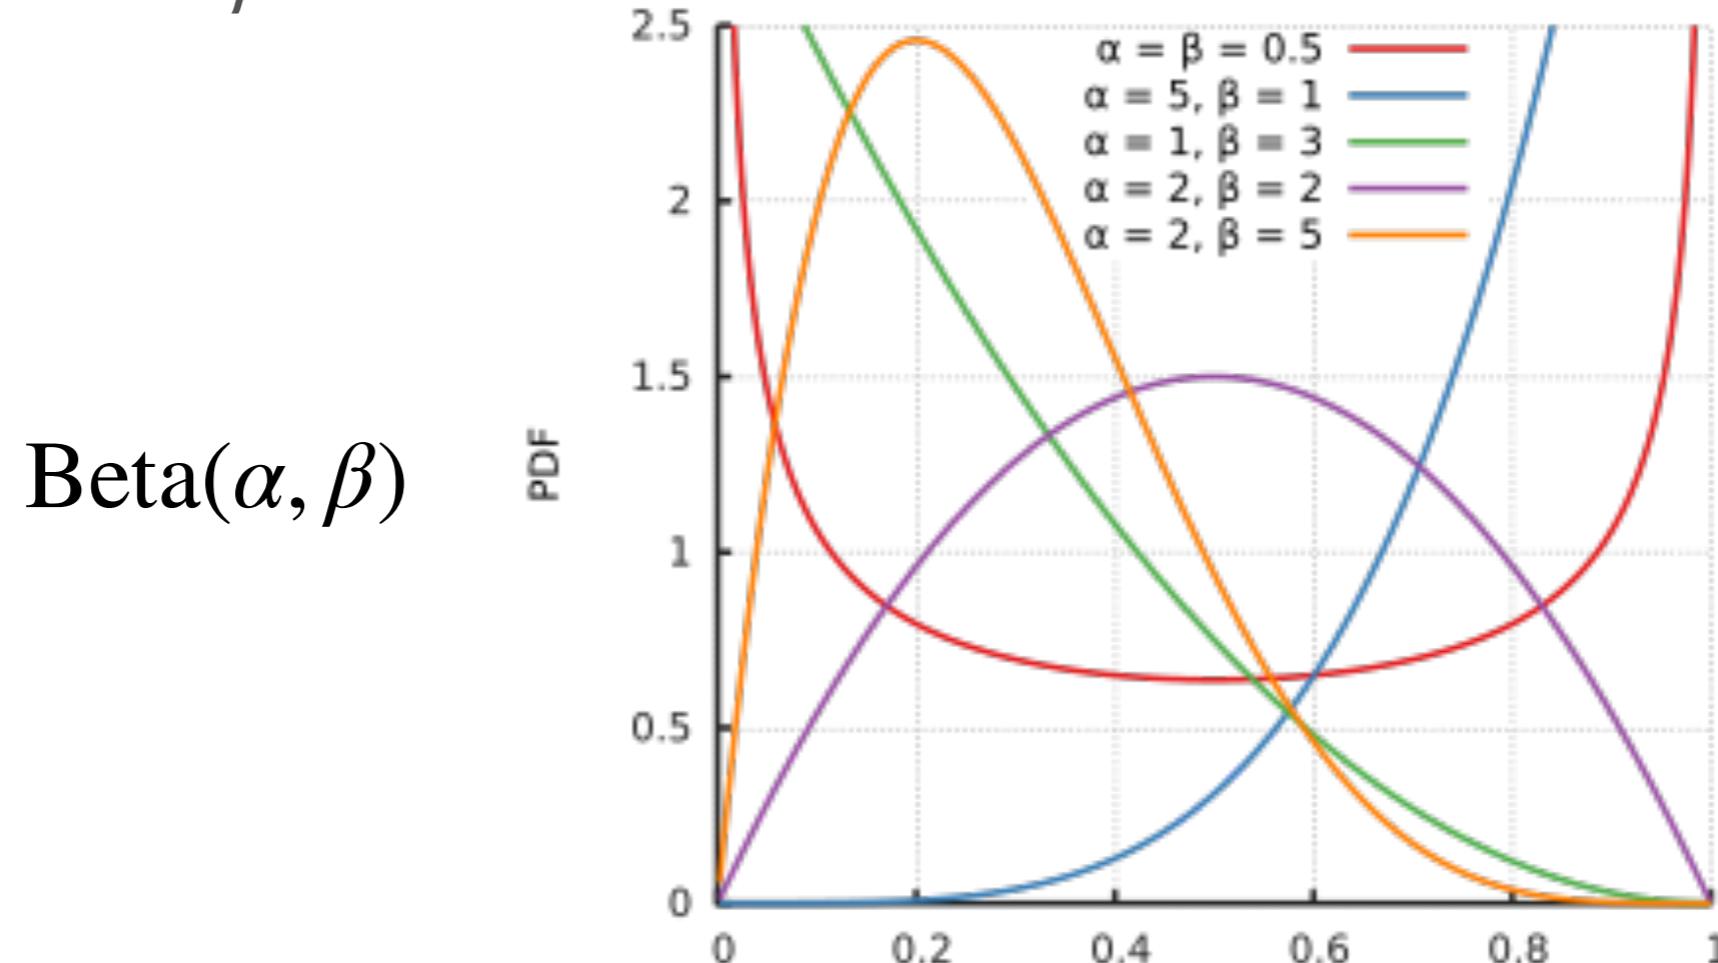
Bernoulli bandits - Prior

Let's consider a Beta distribution **prior** over the mean rewards of the Bernoulli bandits:

$$p(\theta_k) = \frac{\Gamma(\alpha_k + \beta_k)}{\Gamma(\alpha_k)\Gamma(\beta_k)} \theta_k^{\alpha_k-1} (1-\theta_k)^{\beta_k-1} \quad \Gamma(n) = (n-1)!$$

The mean is $\frac{\alpha}{\alpha + \beta}$

The larger the $\alpha + \beta$ the more concentrated the distribution



Bernoulli bandits-Posterior

Let's consider a Beta distribution prior over the mean rewards of the Bernoulli bandits:

$$p(\theta_k) = \frac{\Gamma(\alpha_k + \beta_k)}{\Gamma(\alpha_k)\Gamma(\beta_k)} \theta_k^{\alpha_k-1} (1 - \theta_k)^{\beta_k-1} \quad \Gamma(n) = (n-1)!$$

The posterior is also a Beta! Because beta is conjugate distribution for the Bernoulli distribution.

A closed form solution for the bayesian update, possible only for conjugate distributions!

$$(\alpha_k, \beta_k) \leftarrow \begin{cases} (\alpha_k, \beta_k) & \text{if } x_t \neq k \\ (\alpha_k, \beta_k) + (r_t, 1 - r_t) & \text{if } x_t = k. \end{cases}$$

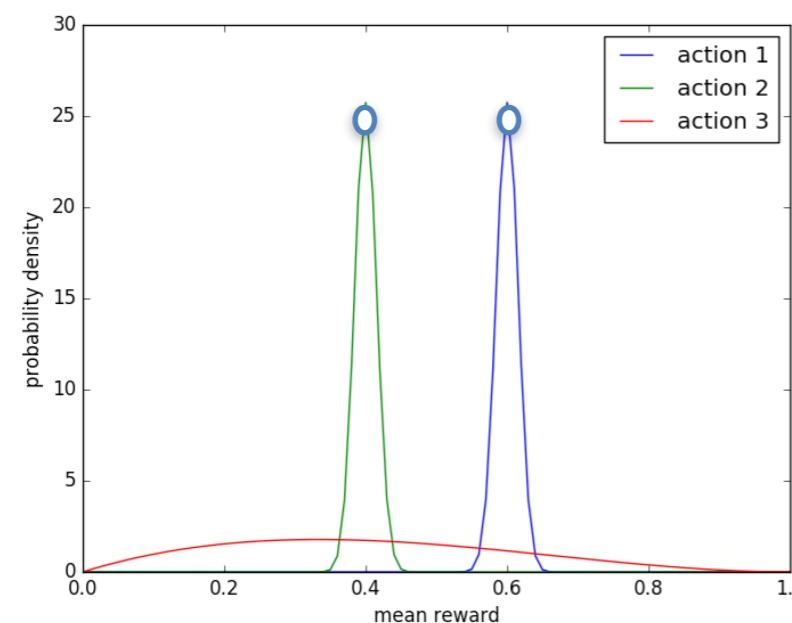
Greedy VS Thompson for Bernoulli bandits

a: success
b: failure

Algorithm 1 BernGreedy(K, α, β)

```

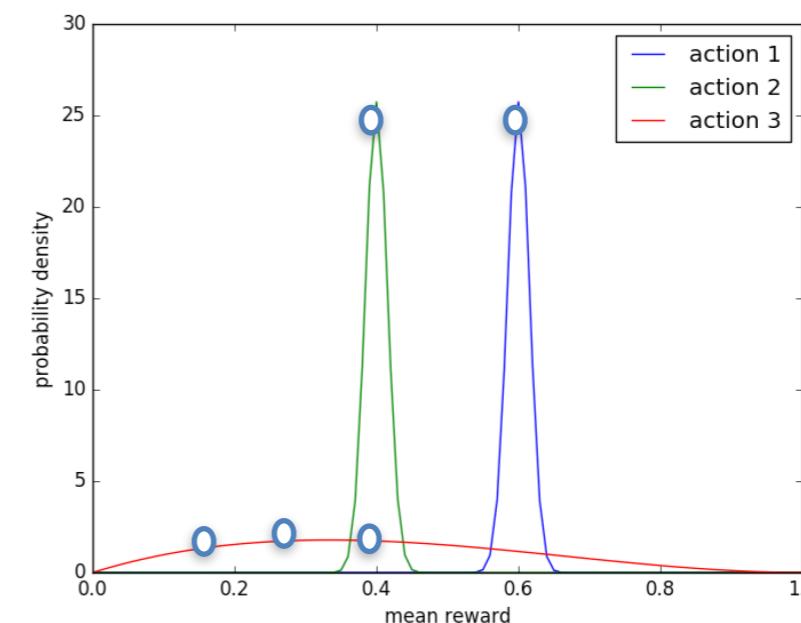
1: for  $t = 1, 2, \dots$  do
2:   #estimate model:
3:   for  $k = 1, \dots, K$  do
4:      $\hat{\theta}_k \leftarrow \alpha_k / (\alpha_k + \beta_k)$ 
5:   end for
6:
7:   #select and apply action:
8:    $x_t \leftarrow \text{argmax}_k \hat{\theta}_k$ 
9:   Apply  $x_t$  and observe  $r_t$ 
10:
11:  #update distribution:
12:   $(\alpha_{x_t}, \beta_{x_t}) \leftarrow (\alpha_{x_t}, \beta_{x_t}) + (r_t, 1 - r_t)$ 
13: end for
```



Algorithm 2 BernThompson(K, α, β)

```

1: for  $t = 1, 2, \dots$  do
2:   #sample model:
3:   for  $k = 1, \dots, K$  do
4:     Sample  $\hat{\theta}_k \sim \text{beta}(\alpha_k, \beta_k)$ 
5:   end for
6:
7:   #select and apply action:
8:    $x_t \leftarrow \text{argmax}_k \hat{\theta}_k$ 
9:   Apply  $x_t$  and observe  $r_t$ 
10:
11:  #update distribution:
12:   $(\alpha_{x_t}, \beta_{x_t}) \leftarrow (\alpha_{x_t}, \beta_{x_t}) + (r_t, 1 - r_t)$ 
13: end for
```



Contextive Bandits (a.k.a Associative Search)

- A contextual bandit is a tuple (A, S, R)
- A is a known set of k actions (or “arms”)
- $\mathcal{S} = \mathbb{P}[s]$ is an unknown distribution over states (or “contexts”)
- $\mathcal{R}_s^a(r) = \mathbb{P}[r | s, a]$ is an unknown probability distribution over rewards
- At each time t
 - Environment generates state $s_t \sim \mathcal{S}$
 - Agent selects action $a_t \in \mathcal{A}$
 - Environment generates reward $r_t \sim \mathcal{R}_{s_t}^{a_t}$
- The goal is to maximize cumulative reward $\sum_{\tau=1}^t r_\tau$

Bayesian Bandits VS A/B Hypothesis testing

- Given two actions, A,B, traditional hypothesis testing would require to gather enough samples to decide whether a preference of A versus B is statistically significant.
- Bayesian bandits adapt the actions over time based on the rewards seen thus far: they accomplish smaller cumulative regret

<https://analytics.googleblog.com/2013/01/multi-armed-bandit-experiments.html>

Real world motivation: Personalized NETFLIX artwork



For a particular title **and a particular user**, we will use the **contextual** multi-armed bandit formulation to decide what image to show per title **per user**

- Actions: uploading an image (available for this movie title) to a user's home screen
- Mean rewards (unknown): the % of NETFLIX users that will click on the title and watch the movie
- Estimated mean rewards: the average click rate (+quality engagement, not clickbait)
- **Context** (s) : user attributes, e.g., language preferences, gender of films she has watched, time and day of the week, etc.

Netflix Artwork: <https://medium.com/netflix-techblog/artwork-personalization-c589f074ad76>

Question: was there a train and test phase on our multi-armed bandit algorithms?

No, the setup we explored was: given a set of K arms, how do we select actions to minimize our cumulative regret.

Q: what would be the learning based equivalent of the multi-armed bandit problem?

A:

- We have a training set of N multi-armed bandit instantiations.
- Each K -armed bandit is one training example.
- The agent gets n number of interactions, and obtains a final reward (-regret).
- The agent learns a policy —mapping from its set of actions taken thus far **and** their outcomes, to a probability over what actions to try next

We will visit this setup in the meta-learning lecture.